From Novel View Synthesis to Generative World Modeling: A Comprehensive Survey

ElephantFlow Research
github.com/elephantflow/ElephantSurvey
Draft v0.3 — March 2026

Abstract

Novel View Synthesis (NVS)—rendering a scene from arbitrary unseen viewpoints given reference observations—has undergone a paradigm shift from classical geometry-based reconstruction toward learning-driven generative approaches. This survey examines NVS not merely as a rendering technique but as the **foundational spatial capability** of Generative World Models: systems that predict visual scene states under arbitrary conditions including viewpoint changes, temporal dynamics, and agent interactions.

We present a unified conceptual framework organising existing methods along a *continuous spectrum*—from static multi-view reconstruction to dynamic, interactive world generation—across three axes: (i) scene representation (implicit, explicit, hybrid), (ii) generative paradigm (per-scene optimisation, feed-forward inference, diffusion-based synthesis), and (iii) world-modelling capability (geometry, dynamics, controllability). We survey more than **180 representative methods** spanning NeRF, 3D Gaussian Splatting, generalizable reconstruction, diffusion-based NVS, and 4-D world modeling, and propose a roadmap toward physically grounded, generalisable world models.

1 Introduction

1.1 The Fundamental Problem

Intelligence requires reasoning about the world beyond what is immediately visible. A self-driving vehicle must anticipate the scene around a blind corner; a robotic arm must predict how a grasped object will appear from the wrist camera after a planned motion; a surgeon must mentally reconstruct anatomy from a 2-D scan. In each case, the core operation is identical: **synthesising a plausible visual state at an unobserved vantage point**.

This is the problem of **Novel View Synthesis (NVS)**. Formally, given a set of reference observations $\{\mathbf{I}_i, \boldsymbol{\pi}_i\}_{i=1}^n$ with known camera poses $\boldsymbol{\pi}_i$, the goal is to render a photo-realistic image $\hat{\mathbf{I}}$ at an arbitrary target pose $\boldsymbol{\pi}^*$, faithfully preserving 3-D geometry, surface appearance, and lighting. The problem is inherently ill-posed when n is small or when the target pose is far from any reference: an infinite family of 3-D scenes is consistent with any finite set of 2-D projections. Resolving this ambiguity requires either dense capture (classical IBR) or a learned prior over the distribution of plausible scenes (modern neural and generative approaches).

While NVS was studied in the image-based rendering era [1–4], the past six years have witnessed an extraordinary paradigm shift: NeRF [17] demonstrated that a shallow MLP can represent complex scenes with view-dependent appearance by differentiably volume-rendering a continuous radiance field; 3DGS [46] showed that an explicit set of anisotropic Gaussians can render at real-time frame

Preprint.

Table 1: Comparison with related surveys. ✓ covered; ∼ partially; – not covered.

Survey	Year	NeRF	3DGS	Generative	Dynamic	Generalise	4D/WM	Unified
Tewari *et al.* [9]	2022	✓	–	∼	∼	∼	–	–
Gao *et al.* [10]	2022	✓	–	–	∼	–	–	–
Zhu *et al.* [11]	2023	∼	–	–	–	–	–	–
Wu *et al.* [12]	2024	–	✓	–	∼	–	–	–
Fei *et al.* [13]	2024	–	✓	∼	∼	–	–	–
Chen *et al.* [14]	2024	∼	∼	–	✓	–	∼	–
Zhang *et al.* [15]	2025	∼	∼	–	–	✓	–	–
He *et al.* [16]	2025	–	✓	∼	∼	–	–	–
Ours	2026	✓	✓	✓	✓	✓	✓	✓

rates; and diffusion-based methods [91, 102, 110] have moved the frontier toward *generalisable, generative* synthesis that requires as few as a single image—or even a text description—as input.

1.2 From NVS to World Modeling

We argue that this progression reflects a deeper conceptual shift: **NVS is becoming the foundational spatial inference engine of Generative World Modeling**.

> **Working Definition (World Model).** A *Generative World Model* is a learned system that can predict the visual (and optionally physical) state of a scene under arbitrary conditions—including viewpoint changes, temporal evolution, and agent interactions—in a manner that is *geometrically consistent*, *physically plausible*, and *generalisable* to scenes and conditions outside the training distribution.

Under this definition, NVS is the special case in which the only varying condition is camera pose. A full world model must additionally handle temporal dynamics (D3), action conditioning (D4), and physical plausibility (D5). Crucially, *spatial consistency enforced by NVS is a necessary prerequisite*: a world model without geometric grounding is merely a video predictor—perceptually plausible in distribution but unreliable for tasks requiring precise spatial reasoning such as robot manipulation or navigation.

Recent systems are actively closing this gap. Genie 2 [125] generates interactive 3-D environments from a single image; Cosmos [126] trains world foundation models at internet scale; and ViewCrafter [110] with StreetCrafter [200] demonstrate that camera-controlled video generation is practical at high resolution. These developments reflect the field's convergence on a unified model that subsumes NVS, video generation, and world simulation.

1.3 Relationship to Prior Surveys

Table 1 compares this work with the most closely related prior surveys. Tewari *et al.* [9] and Gao *et al.* [10] focus on NeRF representations without addressing world modelling. Wu *et al.* [12] and Fei *et al.* [13] cover 3DGS but not generative methods or 4-D. Chen *et al.* [14] cover NVS for autonomous driving but lack a world-modeling lens. To our knowledge, **no existing survey jointly covers all three paradigms (NeRF, 3DGS, generative) through the unifying lens of Generative World Modeling**, nor proposes a continuous-spectrum taxonomy from static reconstruction to 4-D interactive generation.

1.4 The Unified Spectrum

We organise all methods along a *continuous spectrum* defined by three axes (Figure 1):

Axis 1. Input regime: dense multi-view → sparse-view → single image → text/unconditional.

Axis 2. Temporal scope: static scene → deformable/dynamic → 4-D world generation.

Axis 3. Controllability: viewpoint only → viewpoint + time → viewpoint + time + action.

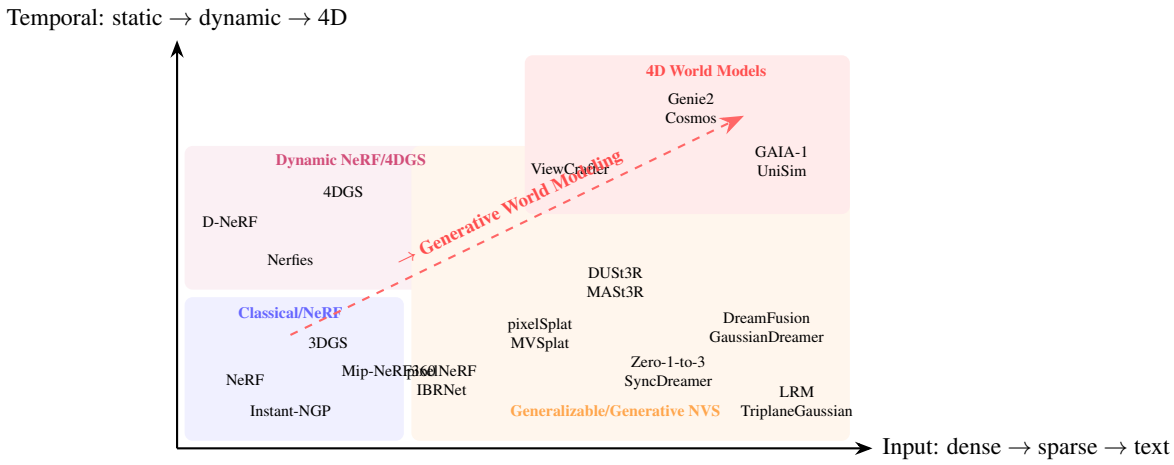

Figure 1: The unified NVS spectrum. The diagonal arrow marks the evolutionary trajectory toward 4-D Generative World Modeling.

1.5 Contributions and Paper Structure

Contributions: (1) A unified continuous-spectrum taxonomy across three axes (input regime, temporal scope, controllability); (2) Cross-paradigm coverage of >180 methods with explicit World Modeling implications in each chapter; (3) A related-survey gap analysis (Table 1) substantiating the novelty of the World Modeling lens; (4) A five-direction open-problems roadmap.

The remainder is organised as follows. §2 establishes background and a unified formulation. §4–6 survey the three representation paradigms (NeRF, 3DGS, feed-forward reconstruction). §7 covers generative NVS with particular depth on camera-controlled diffusion. §8 examines 4-D generation, neural driving simulators, and world foundation models. §9–10 review datasets and evaluation protocols. §11 proposes a research roadmap.

2 Background and Unified Formulation

2.1 Classical Image-Based Rendering

Before the neural era, NVS was approached through *image-based rendering* (IBR) [1–3]. The foundational insight is that images directly encode light transport: sufficiently dense sampling of the plenoptic function [6] permits view interpolation without reasoning about explicit 3-D geometry.

Light Field Rendering [1] and the Lumigraph [5] parameterise the light field on a two-plane slab, enabling photo-realistic view interpolation within the captured light-field volume. View interpolation [3] warps and blends reference images using depth estimates, anticipating modern depth-based view synthesis. Unstructured Lumigraphs and view-dependent texture mapping [4] relaxed capture requirements to unstructured camera sets. Structure from Motion (SfM) [7] and Multi-View Stereo (MVS) [8] established the geometry-estimation pipeline that underlies all pose-conditioned NVS to this day.

The fundamental limitation of classical IBR is the requirement for *dense* capture: interpolation quality degrades rapidly as target viewpoints move beyond the convex hull of the observed camera set. Furthermore, these methods produce no compact scene model—each novel view requires querying the full set of reference images at inference time. The desire for compact, generalisable scene representations drove the development of neural radiance fields.

2.2 Volume Rendering and NeRF

NeRF [17] represents a scene as a continuous function $F_\theta : (\mathbf{x}, \mathbf{d}) \mapsto (\mathbf{c}, \sigma)$, where $\mathbf{x} \in \mathbb{R}^3$ is a 3-D location and $\mathbf{d} \in \mathbb{S}^2$ a viewing direction. The expected colour of a ray $\mathbf{r}(s) = \mathbf{o} + s\mathbf{d}$ is computed

via classical volume rendering [204]:

$$\hat{C}(\mathbf{r}) = \int_{s_n}^{s_f} T(s)\,\sigma(\mathbf{r}(s))\,\mathbf{c}(\mathbf{r}(s),\mathbf{d})\,\mathrm{d}s, \quad T(s) = \exp\left(-\int_{s_n}^{s} \sigma(\mathbf{r}(u))\,\mathrm{d}u\right), \tag{1}$$

where $T(s)$ is the accumulated transmittance. F_θ is a position-encoded MLP [205] trained by minimising $\|\hat{C}(\mathbf{r}) - C(\mathbf{r})\|_2^2$ over all observed rays. Hierarchical stratified sampling decomposes the continuous integral into coarse and fine networks, concentrating samples near the visible surface.

NeRF's key properties—implicit representation, continuous differentiability, and view-dependent appearance via directional encoding—made it transformative. Its central limitations are *slow per-scene optimisation* (hours) and *per-scene specificity* (no generalisation to unseen scenes), both of which the subsequent literature directly addressed.

2.3 Gaussian Splatting

3DGS [46] represents a scene as N anisotropic 3-D Gaussians $\{(\boldsymbol{\mu}_i, \boldsymbol{\Sigma}_i, \alpha_i, \mathbf{c}_i)\}_{i=1}^N$, where $\boldsymbol{\mu}_i$ is the centre, $\boldsymbol{\Sigma}_i$ the full covariance, α_i the opacity, and $\mathbf{c}_i$ a set of spherical-harmonic coefficients encoding view-dependent colour. Novel views are rendered by projecting each Gaussian to a 2-D splat and compositing front-to-back:

$$\hat{C} = \sum_{i \in \mathcal{N}} \mathbf{c}_i\,\alpha_i\,\tilde{G}_i \prod_{j<i} (1 - \alpha_j\,\tilde{G}_j), \tag{2}$$

where $\tilde{G}_i$ denotes the 2-D projection of Gaussian i. Training proceeds by differentiable rasterisation and minimisation of a combined ℓ_1 and SSIM photometric loss. An *adaptive densification* strategy clones under-reconstructed Gaussians and splits over-large ones, while pruning transparent or redundant primitives.

3DGS achieves real-time rendering (>100 fps at 1080p) with training times of tens of minutes—a dramatic improvement over NeRF. Its explicit, editable representation makes it attractive for downstream tasks in scene editing, simulation, and embodied AI. The primary limitations are *higher memory footprint* (millions of Gaussians per scene) and *limited surface quality* compared to SDF-based methods. A large subsequent literature has addressed both limitations.

2.4 Unified NVS and World Modeling Formulation

We formulate NVS as a general function:

$$\mathcal{M}_\theta : \left(\{\mathbf{I}_i, \boldsymbol{\pi}_i\}_{i=1}^n,\ \boldsymbol{\pi}^*\right) \longrightarrow \hat{\mathbf{I}}^*, \quad n \in [1, +\infty). \tag{3}$$

The world-modeling extension augments this with a time index t and a sequence of agent actions $\mathbf{a}_{0:t}$:

$$\mathcal{M}_\theta : \left(\{\mathbf{I}_i, \boldsymbol{\pi}_i\}_{i=1}^n,\ \mathbf{a}_{0:t},\ t,\ \boldsymbol{\pi}^*\right) \longrightarrow \hat{\mathbf{I}}_t^*. \tag{4}$$

Eq. (3) is the special case $t = 0$, $\mathbf{a} = \varnothing$. A complete world model satisfying Eq. (4) must fulfil five desiderata, which we use throughout the survey: **D1** geometric consistency across viewpoints; **D2** generalisation from sparse or single-image input; **D3** dynamic scene modelling; **D4** action conditioning; **D5** physical plausibility.

3 Scene Representation Taxonomy

The choice of scene representation determines the fundamental trade-offs a method faces: rendering speed, reconstruction accuracy, memory footprint, generalisability, and editability. We identify four families, each reflecting a different position on the speed–compactness–quality Pareto frontier.

Implicit representations (MLPs, hash grids) encode the scene as a continuous function $f_\theta(\mathbf{x})$. Their key advantages are compactness and continuous differentiability, which support high-quality reconstruction from dense observations. Their central limitation is rendering speed: each pixel requires integrating along a ray through many network evaluations. NeRF [17] and its successors (§4) are the canonical family.

Explicit representations (voxels, point clouds, Gaussians) store scene attributes at discrete spatial elements, enabling tile-based or splat-based rasterisation that achieves real-time frame rates. The

trade-off is memory: a faithful 3DGS scene may occupy hundreds of megabytes. 3D Gaussian Splatting [46] (§5) is the dominant current representative.

Hybrid representations combine an explicit spatial structure—typically a triplane [88] or multi-resolution grid [23]—with a shallow MLP decoder. They strike a practical balance between the quality of implicit models and the speed of explicit ones, and are the backbone of most large reconstruction models (§6).

Compositional representations decompose the scene into semantic entities (objects, backgrounds, agents) with independent sub-representations. This structural inductive bias enables object-level editing and compositional generation [59, 87], but typically requires supervision or segmentation to separate entities. It is also the natural representation for world models that must reason about independently acting objects.

4 NeRF: Implicit Neural Radiance Fields

4.1 Foundational NeRF Methods

The original NeRF [17] demonstrated that an MLP with positional encoding could faithfully represent complex real scenes with view-dependent appearance from a few dozen posed photographs. Subsequent work systematically addressed its two principal limitations—slow optimisation and poor generalisation to novel scenes.

Accelerating per-scene optimisation. The bottleneck in vanilla NeRF is the MLP query: every rendered pixel requires hundreds of network evaluations along the ray. Instant-NGP [23] replaces the large MLP backbone with a compact multi-resolution *hash grid*, enabling training convergence in under a minute on a single GPU while matching or exceeding prior quality. TensoRF [24] decomposes the radiance field into low-rank CP and VM tensor components, reducing both storage and rendering overhead. Plenoxels [25] eliminates the MLP entirely, using a sparse 3-D voxel grid with trilinearly interpolated spherical harmonics; the resulting model trains in minutes without gradient-intensive network backward passes. DirectVoxGO [26] accelerates convergence via a two-stage coarse-to-fine voxel grid that aggressively skips empty space.

Improving visual quality and handling challenging scenes. Mip-NeRF [19] replaces the point-sampled ray with an *integrated positional encoding* over a conical frustum, substantially reducing aliasing at multiple viewing distances. Mip-NeRF 360 [20] extends this to unbounded outdoor scenes via a non-linear scene parameterisation that contracts the infinite background into a bounded representation. Zip-NeRF [21] unifies the hash grid of Instant-NGP with the anti-aliasing of Mip-NeRF 360, establishing the current quality frontier for per-scene NeRF. NeRF-W [18] introduces per-image latent appearance codes and a transient object model, enabling reconstruction of tourist-photographed landmarks from uncontrolled internet images. Ref-NeRF [22] restructures the directional encoding into a physically-based reflection model that explicitly predicts diffuse colour, specular tint, and surface roughness, dramatically improving the rendering of specular and metallic materials.

Surface reconstruction. While volume rendering naturally produces accurate alpha composites, it does not enforce a well-defined surface, leading to noisy geometry. NeuS [27] reformulates volume rendering using a logistic density derived from a Signed Distance Function (SDF), yielding watertight, high-quality surface reconstruction. VolSDF [28] derives a rigorous analytical mapping from SDF values to volume density, providing a theoretically grounded alternative. BakedSDF [29] combines the surface quality of SDF-based NeRF with real-time rendering by extracting and baking the learned SDF into a high-quality mesh.

Table 2 summarises representative NeRF variants.

4.2 Generalizable NeRF

Per-scene NeRF requires lengthy per-scene optimisation, limiting scalability. *Generalisable* methods learn a shared prior across scenes, enabling fast inference on new scenes.

Table 2: NeRF variants: representative methods and properties.

Method	Venue	Train	Render	Generalise	Key Contribution
NeRF [17]	ECCV'20	Slow	Slow	–	MLP + positional enc.
NeRF-W [18]	CVPR'21	Slow	Slow	–	In-the-wild appearance
Mip-NeRF [19]	ICCV'21	Slow	Slow	–	Anti-aliased cones
Plenoxels [25]	CVPR'22	Fast	Fast	–	No-MLP voxel grid
TensoRF [24]	ECCV'22	Med	Fast	–	Tensor decomp.
DirectVoxGO [26]	CVPR'22	Fast	Fast	–	Two-stage voxel
Instant-NGP [23]	TOG'22	Fast	Fast	–	Hash grid
Mip-NeRF360 [20]	CVPR'22	Slow	Slow	–	Unbounded scenes
Ref-NeRF [22]	CVPR'22	Slow	Slow	–	Specular reflection
NeuS [27]	NeurIPS'21	Slow	Slow	–	SDF-based surface
VolSDF [28]	NeurIPS'21	Slow	Slow	–	SDF-density bridge
Zip-NeRF [21]	ICCV'23	Fast	Fast	–	Hash + anti-alias
BakedSDF [29]	SIGGRAPH'23	Med	Fast	–	Mesh extraction

pixelNeRF [30] conditions the NeRF MLP on pixel-aligned CNN features extracted from reference views, enabling single-image reconstruction. IBRNet [31] uses a transformer to aggregate features from multiple reference views at query points along the ray. MVSNeRF [32] builds a cost volume from stereo for geometry-aware feature extraction. GNT [33] dispenses with the NeRF MLP entirely, replacing it with a two-stage transformer that first encodes reference views into a view-independent representation and then renders the target view via cross-attention, learning geometric reasoning as an emergent property of attention rather than an explicit inductive bias. SRT [34] pursues a similar philosophy but uses a set-latent architecture, encoding an arbitrary number of views into a fixed-size scene representation without pose assumptions. NerfDiff [35] augments generalizable NeRF with diffusion priors for single-image synthesis. ZeroRF [36] achieves sparse-view 360° reconstruction via per-scene optimisation with zero pretraining, leveraging factorised feature grids.

Few-shot and test-time methods. FrameNeRF [162] presents a simple and efficient framework for few-shot NVS by exploiting temporal video frames as dense pseudo views. Recollection from Pensieve [173] learns NVS from uncalibrated internet videos without posed supervision. Test-Time Completion [195] repurposes natural video completion as a zero-shot NVS prior at test time. Mixed Neural Voxels [174] enables fast multi-view video synthesis by decomposing scenes into static and dynamic voxel grids.

4.3 Dynamic NeRF

Modelling non-rigid dynamic scenes is a central challenge for world modeling.

D-NeRF [37] introduces a deformation network mapping each spacetime point to a canonical static space. Nerfies [39] and HyperNeRF [40] model complex non-rigid deformations via latent codes and higher-dimensional ambient spaces. Neural Scene Flow Fields (NSFF) [38] models dynamic scenes with scene flow representation. Non-Rigid NeRF [196] reconstructs and synthesises novel views of dynamic scenes from monocular video with explicit ray bending. DyNeRF [41] extends NeRF to monocular video via a time-conditioned model. TiNeuVox [42] uses voxel features to accelerate dynamic NeRF. K-Planes [43] decomposes 4-D space-time into six feature planes—three spatial (XY, XZ, YZ) and three spatial-temporal (XT, YT, ZT)—each encoded as a learned 2-D grid, making the architecture highly memory-efficient and interpretable. HexPlane [44] independently proposes the same factorisation, demonstrating its generality; together these works establish plane-based 4-D decomposition as the dominant acceleration paradigm for dynamic NeRF. NeRFPlayer [45] separately encodes static, deforming and newly appearing regions. Soup-of-Planes [183] enables fast view synthesis of casual monocular videos by fitting a set of planar scenes. Pseudo-Generalised Dynamic View [184] synthesises dynamic views from monocular video by combining static and dynamic video inpainting. Decoupled Dynamic Video [179] disentangles foreground dynamics and background statics for monocular dynamic view synthesis.

Table 3 summarises dynamic NeRF methods.

Table 3: Dynamic NeRF methods.

Method	Venue	Multi-view	Monocular	Key Idea
D-NeRF [37]	CVPR'21	✓	–	Deformation network
NSFF [38]	CVPR'21	✓	∼	Scene flow field
Non-Rigid NeRF [196]	ICCV'21	–	✓	Ray bending
Nerfies [39]	ICCV'21	–	✓	Latent deformation
HyperNeRF [40]	TOG'21	–	✓	Ambient dimension
DyNeRF [41]	CVPR'22	✓	–	Multi-view video NeRF
TiNeuVox [42]	SIGGRAPH'22	✓	–	Voxel dynamic accel.
K-Planes [43]	CVPR'23	✓	∼	4-D plane decomp.
HexPlane [44]	CVPR'23	✓	∼	Hexagonal planes
NeRFPlayer [45]	TOG'23	✓	–	Region decomp.
Soup-of-Planes [183]	arXiv'23	–	✓	Planar scene fitting
Pseudo-GDV [184]	CVPR'24	–	✓	Static+dyn. inpaint

5 3D Gaussian Splatting

5.1 Foundational 3DGS

3DGS [46] achieved real-time ($>$100 fps at 1080p) NVS with training times of 30–45 minutes on commodity hardware, surpassing NeRF-quality baselines on standard benchmarks. Its three key innovations are: (i) anisotropic Gaussian primitives that align naturally with planar scene structures; (ii) tile-based rasterisation that enables massively parallel GPU rendering; (iii) adaptive densification that automatically allocates capacity to under-reconstructed regions. The resulting representation is fully explicit and editable, opening new directions in scene editing, simulation, and avatar rendering.

Quality and anti-aliasing. Mip-Splatting [47] identifies aliasing in vanilla 3DGS as arising from the approximation of 3-D Gaussian convolution during 2-D projection. It introduces a 3-D smoothing filter derived from the sampling footprint, enabling accurate multi-scale rendering and substantially reducing floaters at large distances. AbsGS [48] revisits the densification criterion: replacing the mean gradient with an absolute-value measure improves coverage in thin structures and reduces needle-like artifacts near object boundaries.

Compression and memory efficiency. A single 3DGS scene may require 300MB–1GB of storage, limiting deployment. Scaffold-GS [49] introduces a sparse set of anchor primitives from which neural networks predict neighbouring Gaussian attributes; this structured representation reduces storage by 3–5$\times$ with comparable quality. LightGaussian [50] achieves 15$\times$ compression via global Gaussian pruning based on visibility contribution and subsequent vector quantisation of attributes. EAGLES [52] applies residual vector quantisation for fine-grained attribute compression. HAC [53] proposes a hash-grid-based context model that predicts per-Gaussian attribute entropy, enabling learned arithmetic coding for extreme compression ratios.

Surface reconstruction. SuGaR [54] aligns Gaussians to the scene surface by adding a regularisation term that encourages Gaussians to become flat discs at the surface, followed by mesh extraction and joint Gaussian–mesh refinement. 2DGS [55] replaces 3-D Gaussians with 2-D disc primitives, which by construction lie on surfaces and produce accurate normal maps. Gaussian Opacity Fields (GOF) [56] fuses Gaussian Splatting with a continuous opacity field for topologically correct surface extraction in unbounded scenes.

Scene editing and compositionality. The explicit Gaussian representation is particularly amenable to editing: individual primitives can be identified, moved, duplicated, or removed. GaussianEditor [57] harnesses a large vision-language model to select semantically relevant Gaussians from a text description and optimises their appearance with a diffusion prior. Gaussian Grouping [58] augments each Gaussian with a learned identity field, trained with SAM [210] masks from rendered views, enabling instance-level part decomposition and editing. GALA3D [59] uses an LLM to generate a 3-D layout from text and initialises Gaussians per-object, enabling compositional scene generation.

Table 4 summarises key 3DGS variants.

Table 4: 3DGS variants.

Method	Venue	RT	Dyn.	Surf.	Comp.	Key Idea
3DGS [46]	TOG'23	✓	–	–	–	Gaussian splatting
Mip-Splatting [47]	CVPR'24	✓	–	–	–	Anti-aliased splats
AbsGS [48]	NeurIPS'24	✓	–	–	–	Absolute grad. dens.
Scaffold-GS [49]	CVPR'24	✓	–	–	✓	Anchor structure
LightGaussian [50]	NeurIPS'23	✓	–	–	✓	Pruning + VQ
Compact3D [51]	arXiv'24	✓	–	–	✓	k-means codebook
EAGLES [52]	ECCV'24	✓	–	–	✓	Residual VQ
HAC [53]	ECCV'24	✓	–	–	✓	Hash context model
SuGaR [54]	CVPR'24	✓	–	✓	–	Joint mesh+GS
2DGS [55]	SIGGRAPH'24	✓	–	✓	–	2-D disk Gaussians
GOF [56]	arXiv'24	✓	–	✓	–	Opacity field
GaussianEditor [57]	CVPR'24	✓	–	–	–	LLM-guided editing
Gaussian Grouping [58]	ECCV'24	✓	–	–	–	Identity field
GALA3D [59]	arXiv'24	✓	–	–	–	Compositional gen.

Table 5: Feed-forward generalizable Gaussian Splatting methods.

Method	Venue	Input views	Unposed	Key Idea
pixelSplat [66]	CVPR'24	2	–	Stereo epipolar feat.
MVSplat [67]	ECCV'24	2–6	–	Cost volume prior
TranSplat [68]	arXiv'24	2+	–	Transformer MV
FSGS [69]	CVPR'24	3	–	Gaussian unpooling
DepthSplat [70]	arXiv'24	2+	–	Joint depth+GS
Splatt3R [71]	arXiv'24	2	✓	DUSt3R backbone
TriplaneGaussian [72]	CVPR'24	1	–	Triplane+Gaussian

5.2 Dynamic 3DGS

4DGS [60] extends Gaussians into 4-D space-time with a temporal radiance field. SpacetimeGaussians [61] assigns polynomial time trajectories to each Gaussian. Deformable3DGS [62] applies a deformation network to canonical Gaussians. SC-GS [63] uses sparse control points for monocular dynamic scenes. Gaussian-Flow [64] models Gaussian dynamics via polynomial motion fields in Fourier space. Shape-of-Motion [65] reconstructs per-Gaussian 3-D motion from monocular videos. Dynamic 3D Gaussians [175] tracks scene dynamics through persistent Gaussian views, jointly optimising NVS and tracking. GauFRe [191] uses Gaussian deformation fields with flow regularisation for real-time dynamic NVS. MoVieS [192] generates motion-aware 4-D dynamic views in approximately one second via efficient Gaussian motion estimation. Self-Calibrating 4D [176] performs 4-D NVS from monocular videos by jointly estimating camera parameters and Gaussian deformations. Reconstruct-Inpaint-Finetune [177] addresses dynamic monocular NVS via a three-stage pipeline of reconstruction, inpainting, and test-time refinement. EVA-Gaussian [185] achieves real-time human NVS under diverse multi-view settings. Articulated Gaussian [194] enables template-free articulated Gaussian Splatting for real-time reposable dynamic view synthesis.

5.3 Feed-Forward Generalizable 3DGS

pixelSplat [66] predicts Gaussian primitives from stereo image pairs in a single forward pass. MVSplat [67] adds cost-volume priors for better geometry. TranSplat [68] introduces a transformer architecture for multi-view Gaussian prediction. Splatt3R [71] builds on DUSt3R [73] for unposed image pairs. FSGS [69] enables few-shot Gaussian Splatting via Gaussian unpooling and monocular depth supervision. DepthSplat [70] jointly predicts depth and Gaussians for generalizable NVS. TriplaneGaussian [72] combines triplane and Gaussian representations in a transformer-based single-view reconstruction model.

Table 5 summarises feed-forward generalizable 3DGS.

6 Generalisable Feed-Forward 3D Reconstruction

6.1 Pose-Free and Foundation Models

A new generation of methods aims for *pose-free* generalizable reconstruction. DUSt3R [73] predicts dense pointmaps from arbitrary image pairs without camera calibration, then aligns them via global optimisation. MASt3R [74] adds matching capabilities to DUSt3R, enabling accurate localisation. MonST3R [75] extends DUSt3R to dynamic video. VGGT [76] is a visual geometry grounded deep structured tokens model for feed-forward 3-D structure prediction.

6.2 Large Reconstruction Models (LRM)

LRM [77] trains a large transformer on objaverse, enabling single-image-to-3D in under 5 seconds. Instant3D [78] generates NeRF from a single image via a diffusion-conditioned LRM. OpenLRM [79] is the open-source counterpart. LGM [80] learns multi-view Gaussian generation for fast single-image 3-D. CRM [81] uses a convolutional reconstruction model.

6.3 Multi-View Stereo in the Neural Era

Classical MVS has been reinvigorated by deep learning. MVSNet [82] introduced cost-volume-based deep MVS. IterMVS [83] and Uni-MVSNet [84] improved accuracy and generalisation. These methods provide the geometry backbone for downstream NVS.

7 Generative Paradigms for Novel View Synthesis

The emergence of deep generative models—first GANs and then diffusion models—has fundamentally altered the landscape of NVS. Where reconstruction-based methods infer a scene representation from observations and render it deterministically, *generative* methods learn a distribution over plausible visual completions given partial observations. This distinction is not merely architectural: it represents a shift from *interpolation* (rendering within the observed convex hull of viewpoints) to *extrapolation* (hallucinating plausible content beyond it). Generative methods are therefore the primary vehicle for the one-image-to-3D, text-to-3D, and single-image world-generation capabilities that bring NVS into the world-modeling regime.

7.1 3-D-Aware GANs: Learning Geometry from 2-D Supervision

The first wave of generative NVS methods embedded radiance fields inside GAN generators, enabling implicit 3-D supervision from unposed 2-D image collections—a paradigm that proved influential precisely because it required no 3-D ground truth.

GRAF [85] pioneered this direction by conditioning a NeRF generator on a latent shape-appearance code, training with a 2-D discriminator. pi-GAN [86] improved the NeRF backbone with periodic activations (SIREN [209]), achieving substantially better multi-view consistency. GIRAFFE [87] introduced scene compositionality by combining foreground and background neural feature fields, enabling object-level editing without explicit segmentation. EG3D [88] is the landmark method of this era: its *triplane* representation—three axis-aligned 2-D feature maps decoded by a shallow MLP—reduces rendering cost by orders of magnitude and admits a super-resolution upsampler, enabling 512×512 training at practical throughput. StyleNeRF [89] integrates the NeRF volume renderer into a StyleGAN [208] backbone, inheriting its powerful style-based latent structure for fine-grained appearance control.

A critical limitation of GAN-based approaches is *mode coverage*: the discriminator objective encourages sharp, realistic individual samples but does not explicitly minimise distributional divergence over all plausible views, making it prone to mode dropping for complex, asymmetric scenes. This motivated the transition toward diffusion-based approaches.

7.2 Diffusion Models for Camera-Controlled NVS

Diffusion models [90, 206] have become the dominant generative paradigm for NVS due to three decisive advantages: (i) training stability via a simple denoising objective; (ii) access to powerful

image priors from large-scale pretraining; (iii) natural integration of arbitrary conditioning signals (pose, depth, text, reference image) through classifier-free guidance.

7.2.1 Pose-Conditioned Single-Image NVS

The most direct application is *single-image NVS*: given one reference photograph and a target camera pose π^*, generate a plausible rendering at that viewpoint. This is under-constrained without a prior—an infinite number of 3-D scenes are consistent with a single 2-D observation—and diffusion models provide exactly the distributional prior needed.

Zero-1-to-3 and successors. Zero-1-to-3 [91] established the canonical formulation: fine-tune a latent diffusion model (Stable Diffusion [90]) on multi-view renders from Objaverse [146], conditioning on both the reference image and the relative pose $(\Delta\phi, \Delta\theta, \Delta r)$ encoded as a CLIP-space vector. The model learns to "imagine" how an object would look from any viewpoint, without constructing an explicit 3-D representation. Zero123++ [94] strengthens multi-view consistency by predicting a fixed grid of six canonical views simultaneously, reducing inter-view inconsistency from independent sampling. One-2-3-45 [93] pipelines Zero-1-to-3 outputs into a cost-volume-based MVS reconstruction [82], converting the diffusion output to a textured mesh within 45 seconds.

Synchronised multi-view denoising. A key challenge with independently conditioned diffusion is that each generated view is sampled from a different realisation of the latent distribution, producing inter-view geometric contradictions. SyncDreamer [92] addresses this by coupling the denoising of multiple target views through shared 3-D volume attention at each diffusion step, enforcing consistent geometry throughout the synthesis process. MVDiffusion [95] constructs a correspondence-aware attention mechanism that explicitly queries matching pixels across views, enabling holistic panorama generation from a set of overlapping perspective images. Wonder3D [96] takes a two-domain approach: a cross-domain diffusion model simultaneously denoises an RGB and a surface-normal image, with the normal output serving as high-quality geometric supervision for downstream mesh reconstruction.

Scalable and pose-free generation. EscherNet [97] generalises pose conditioning via a novel *camera positional encoding* (CaPE), enabling generation of an arbitrary number of views at any resolution without per-angle fine-tuning. ViVid-1-to-3 [169] recasts NVS as a video generation problem: the temporal axis of a pre-trained video diffusion model is repurposed to model the viewpoint axis, using the model's temporal consistency as a proxy for multi-view coherence. GeoNVS [168] addresses geometric fidelity by conditioning a video diffusion model on a rendered point cloud at the target pose, grounding the generative hallucination in explicit 3-D geometry.

7.2.2 Multi-View Diffusion and 3-D-Consistent Generation

Single-view methods generate views independently or with limited cross-view coupling. A more ambitious goal is to train a diffusion model that *natively* produces multi-view consistent outputs—effectively encoding 3-D structure implicitly in the model weights.

MVDream [101] achieves this by replacing the temporal self-attention in a video diffusion transformer with a 3-D-aware *multi-view attention* block that attends across all views simultaneously. Training on rendered 3-D assets (Objaverse) with known camera poses teaches the model to share structural information across views. The resulting model can be used both as a standalone multi-view generator and as a geometry-consistent score function for SDS-based text-to-3D.

GeNVS [100] takes a hybrid approach: a NeRF encoder constructs a 3-D feature volume from the conditioning images, and the diffusion decoder renders a target view by volumetrically querying this feature space. This architecture provides hard geometric grounding while preserving the expressiveness of the diffusion prior for high-frequency appearance details.

7.2.3 Camera Trajectory Conditioning and Video Diffusion NVS

A particularly important frontier is conditioning diffusion models on *continuous camera trajectories*, enabling the generation of smooth, temporally coherent novel-view sequences—the direct bridge to world modeling.

Camera trajectory as a control signal. ViewCrafter [110] conditions a video diffusion model (CogVideoX) on point-cloud renderings along a target camera path, using the rendered point cloud as a structural scaffold that prevents temporal drift while the diffusion model fills in appearance and fine geometry. Free3D [99] injects camera pose into each frame's cross-attention query, enabling a single video diffusion model to follow arbitrary camera trajectories from a single reference image. Novel View Extrapolation [172] specifically targets *large-angle extrapolation*—the regime where geometric hallucination is most necessary—using video diffusion priors to maintain perceptual plausibility when the reference content provides little direct evidence.

Training-free and zero-shot camera control. NVS-Solver [193] demonstrates that an off-the-shelf video diffusion model already encodes sufficient geometric knowledge to solve NVS in a zero-shot manner: by formulating view synthesis as a constrained diffusion process with known-region guidance, it achieves competitive performance without any pose-specific fine-tuning. DDIM Inversion for NVS [163] exploits the deterministic reconstruction property of DDIM [207]: inverting the reference image to its latent code and re-decoding with a modified pose conditioning produces structurally faithful novel views, again without additional training.

Geometric conditioning for high fidelity. Projected Representation Conditioning [166] proposes projecting 3-D features from the reference view onto the target-view image plane, providing an explicit geometric "warp" prior to the diffusion decoder. This projected feature conditioning substantially reduces hallucination artifacts in regions visible from the target but occluded in the reference. A Single Image and Multimodality [167] augments pose conditioning with depth, surface normal, and semantic maps, each derived from monocular estimation, to enrich the conditioning signal across multiple complementary geometric cues.

7.2.4 Reconstruction–Diffusion Hybrids

A complementary design philosophy *separates* the diffusion and reconstruction stages: diffusion fills in missing views and hallucinations, while a deterministic reconstruction pipeline (NeRF or 3DGS) enforces 3-D consistency.

Cat3D [108] exemplifies this pattern: given sparse reference views ($n \geq 1$), a multi-view diffusion model generates a dense set of synthetic pseudo-views, which are then fed to a standard 3DGS reconstruction pipeline. Because the diffusion model is fine-tuned jointly with the reconstruction loss, the generated views are explicitly optimised for downstream reconstruction fidelity. ReconFusion [109] integrates the diffusion prior more tightly: it conditions a pixel-diffusion model on pixelNeRF [30] feature queries, acting as a learned regulariser that fills in uncertain regions during sparse-view NeRF optimisation. ZeroNVS [98] handles the challenging case of in-the-wild single images (non-object-centric, background present) by training on scene-level data and conditioning on the camera intrinsics alongside pose, enabling 360° generation from a single casual photograph. NVComposer [170] addresses the case where multiple reference views are available but unposed: a feature composition network aggregates information from all reference images before passing to the diffusion model, improving quality without requiring explicit pose estimation.

7.3 Score Distillation: Diffusion as a 3-D Optimisation Oracle

Score Distillation Sampling (SDS) [102] opened a separate pathway: instead of directly generating novel views, use a pre-trained 2-D diffusion model as a differentiable *score function* to optimise a 3-D representation. The core identity is:

$$\nabla_\theta \mathcal{L}_{\text{SDS}} = \mathbb{E}_{t,\boldsymbol{\epsilon}}\left[w(t)(\hat{\boldsymbol{\epsilon}}_\phi(\mathbf{x}_t; y, t) - \boldsymbol{\epsilon}) \, \frac{\partial \mathbf{x}}{\partial \theta} \right], \tag{5}$$

where θ parameterises the 3-D representation, $\mathbf{x}$ is a rendered image, y is a text prompt, and $\hat{\boldsymbol{\epsilon}}_\phi$ is the diffusion model's noise prediction.

DreamFusion [102] first demonstrated that SDS can synthesise coherent 3-D objects from text descriptions. Magic3D [103] introduces a coarse-to-fine strategy that first optimises a coarse NeRF then refines a differentiable mesh, achieving $8\times$ higher resolution. ProlificDreamer [104] identifies the "over-saturation" artifact of SDS as a consequence of mode-seeking behaviour, proposing *Variational Score Distillation* (VSD) which models the 3-D parameter distribution as a particle set and applies

Stein variational gradient descent for diverse, high-quality generation. Fantasia3D [105] disentangles geometry (optimised via SDS on a DMTet mesh) from appearance (a physically-based BRDF), enabling photorealistic materials and independent editing of shape and texture. GaussianDreamer [106] replaces NeRF with 3DGS as the SDS target, leveraging Gaussian densification for fine detail and achieving real-time rendering of the final asset. LucidDreamer [107] addresses SDS instability with an interval-based score matching scheme that reduces gradient variance.

7.4 Dedicated Multi-View Diffusion Architectures

The methods above predominantly adapt existing video or image diffusion models. A complementary line of work designs architectures *from scratch* for multi-view consistency.

MVDream [101] replaces temporal attention with multi-view attention and trains on Objaverse with ground-truth camera poses, producing natively 3-D-consistent outputs. ImageDream [212] extends MVDream to image-conditioned generation, enabling single-image-to-3D with strong semantic fidelity. Zero123-XL [211] scales the Zero-1-to-3 paradigm to larger model capacities and more diverse training data. Direct3D [213] trains a native 3-D latent diffusion model on a triplane latent space, enabling direct 3-D generation without a two-stage pipeline.

Table 6 provides a structured comparison of generative NVS methods across the four paradigms above.

Implications for World Modeling. Diffusion-based NVS represents a critical advance along World Modeling Axis 2 (sparse/single-image generalisation, **D2**) and provides the generative prior essential for completing unobserved scene regions (Axis 1, geometric consistency **D1**). Camera trajectory conditioning (§7.2) directly bridges into Axis 3 (temporal dynamics), as a camera trajectory can be viewed as a special case of a time-ordered action sequence. The remaining gap is *physical action conditioning* (**D4, D5**): existing diffusion NVS methods condition on camera parameters but not on object-level physical interactions, which remains an open problem addressed in §11.

8 Towards Generative World Modeling

8.1 Dynamic Scenes and 4-D World Representation

The step from static NVS to world modeling hinges on modelling *temporal dynamics*: how does a scene evolve over time, and can the model generate novel views at arbitrary future timestamps? The dynamic NeRF and 4DGS methods in §4–5 address this for *observed* scenes given dense multi-view video. A harder and more practically relevant regime is *generalizable* 4-D generation from sparse or single-image input—precisely the regime inhabited by world models.

Video diffusion models have emerged as the most promising backbone. ViewCrafter [110] augments a video diffusion model with point-cloud renderings along the target camera trajectory, providing a geometric scaffold that constrains the generation while the diffusion prior fills in appearance and fine detail. CAT4D [199] generalises the Cat3D framework to the temporal dimension, jointly modelling viewpoint and time with a multi-view video diffusion model, then lifting the output to an explicit 4-D representation. 4DiM [111] takes a more direct approach, training a 4-D diffusion model that simultaneously denoises a set of frames at varying timestamps and viewpoints, learning a joint spatiotemporal prior. SV3D [198] and SV4D [112] extend Stable Video Diffusion to 3-D and 4-D object generation respectively, conditioning on orbital camera paths to ensure multi-view consistency. Consistent4D [113] enforces temporal consistency across a 4-D cascade, resolving inter-frame Gaussian drift. DreamGaussian4D [114] shows that SDS can be applied to 4DGS, enabling text-driven 4-D generation without any video training data.

A complementary line of work addresses the *monocular* case. Diffusion Priors for Dynamic NVS [180] treats the diffusion model as a frame-level regulariser during per-scene 4-D optimisation from a single monocular video, substantially improving novel-view quality in under-constrained regions.

Table 6: Generative NVS methods. **Backbone**: LDM=latent diffusion, VDM=video diffusion, GAN=generative adversarial. **Ctrl**: pose conditioning mechanism. **3D Out**: whether the method produces an explicit 3-D representation.

Method	Venue	Backbone	Ctrl	3D Cons.	3D Out	Key Contribution
GAN-based						
GRAF [85]	NeurIPS'20	NeRF-GAN	Latent	$\sim$	NeRF	Radiance field GAN
pi-GAN [86]	CVPR'21	NeRF-GAN	Latent	$\sim$	NeRF	Periodic implicit
GIRAFFE [87]	CVPR'21	Feat. Field	Latent	$\sim$	Neural	Compositional
EG3D [88]	CVPR'22	Triplane	Latent	✓	Tri.	Triplane+SR
StyleNeRF [89]	ICLR'22	StyleGAN	Latent	$\sim$	NeRF	Style-based NeRF
Pose-conditioned single-image diffusion						
Zero-1-to-3 [91]	ICCV'23	LDM	Δpose	$\sim$	–	Pose-cond LDM
Zero123++ [94]	arXiv'23	LDM	Δpose	$\sim$	–	6-view grid
One-2-3-45 [93]	NeurIPS'23	LDM+MVS	Δpose	✓	Mesh	Diff$\rightarrow$MVS
SyncDreamer [92]	ICLR'24	LDM	Δpose	✓	–	Volume sync
Wonder3D [96]	CVPR'24	LDM	Δpose	✓	–	RGB+Normal
EscherNet [97]	CVPR'24	LDM	CaPE	✓	–	Arbitrary #views
ViVid-1-to-3 [169]	arXiv'24	VDM	Δpose	$\sim$	–	Video=NVS axis
Camera trajectory / video diffusion						
Free3D [99]	ICLR'24	VDM	Cam embed	$\sim$	–	Per-frame pose
ViewCrafter [110]	arXiv'24	VDM	Ptcloud	✓	–	Scaffold+VDM
NVS-Solver [193]	arXiv'24	VDM	Zero-shot	$\sim$	–	Guided diffusion
DDIM Inversion [163]	arXiv'24	LDM	Inversion	$\sim$	–	Training-free
GeoNVS [168]	arXiv'24	VDM	Ptcloud	✓	–	Geo-grounded VDM
Proj. Repr. Cond. [166]	arXiv'24	LDM	3D proj.	✓	–	Feature warp
Novel View Extrap. [172]	arXiv'24	VDM	Cam traj.	$\sim$	–	Large-angle extrap.
Reconstruction–diffusion hybrids						
GeNVS [100]	ICLR'23	LDM	NeRF feat	✓	NeRF	NeRF-cond diff.
ZeroNVS [98]	CVPR'24	LDM	Pose+int.	$\sim$	–	In-the-wild
ReconFusion [109]	CVPR'24	LDM	pixelNeRF	✓	NeRF	Diff reg.
Cat3D [108]	NeurIPS'24	MV-Diff	Δpose	✓	3DGS	MV gen+recon
NVComposer [170]	arXiv'24	LDM	Feature	✓	–	Unposed multi-ref
Score distillation (text-to-3D)						
DreamFusion [102]	ICLR'23	LDM+NeRF	SDS	$\sim$	NeRF	SDS
Magic3D [103]	CVPR'23	LDM+Mesh	SDS	$\sim$	Mesh	C2F SDS
ProlificDreamer [104]	NeurIPS'23	LDM+NeRF	VSD	$\sim$	NeRF	VSD
Fantasia3D [105]	ICCV'23	LDM+Mesh	SDS	$\sim$	BRDF	Disentangled
GaussianDreamer [106]	CVPR'24	LDM+3DGS	SDS	$\sim$	3DGS	SDS+3DGS

8.2 Neural Scene Simulators for Autonomous Driving

Autonomous driving provides a uniquely demanding test-bed for NVS-as-world-model: scenes are large-scale, unbounded, and contain independently moving actors (vehicles, pedestrians) that must be rendered consistently from closed-loop ego-motion trajectories. Three architectural strategies have emerged.

NeRF-based compositional simulators. MARS [119] separates the scene into a static NeRF background and per-agent NeRF foregrounds, allowing independent actor manipulation. UniSim [118] extends this to a closed-loop sensor simulator that supports controllable scene editing—removing actors, changing weather, inserting novel agents—with photorealistic rendering at the ego sensor. EmerNeRF [121] learns the dynamic-static decomposition in a self-supervised manner by decomposing scene flow from multi-camera video, without per-actor instance labels.

3DGS-based driving reconstruction. Street Gaussians [122] extends 3DGS to urban scenes by representing tracked foreground objects with separate Gaussian models and compositing them over a background Gaussian field. PVG [123] models Gaussian motion with periodic vibration functions that efficiently encode slow, cyclic motions common in urban traffic. OmniRe [120] unifies background, foreground, and sky in a single Gaussian framework, establishing a strong all-in-one baseline.

Table 7: World modeling and neural scene simulation methods.

Method	Year	3D Aware	Video	Action	Domain	Key Feature
GAIA-1 [117]	2023	∼	✓	✓	Driving	Seq. token WM
UniSim [118]	2023	✓	✓	✓	Driving	Sensor sim.
MARS [119]	2023	✓	✓	∼	Driving	Compositional
Genie [124]	2024	–	✓	✓	Games	Latent action
Genie 2 [125]	2024	✓	✓	✓	Open-world	Interactive 3D
WonderWorld [129]	2024	✓	–	–	Open-world	Instant 3DGS
ViewCrafter [110]	2024	✓	✓	–	Open-world	Cam-cond video
Cosmos [126]	2025	✓	✓	✓	Physical	WFM training
Street Gaussians [122]	2024	✓	✓	∼	Driving	3DGS driving
OmniRe [120]	2024	✓	✓	✓	Driving	All-in-one

Diffusion-based street generation. StreetCrafter [200] uses a point-cloud-conditioned video diffusion model with controllable camera motion, enabling photorealistic street-view synthesis for data augmentation. Sat2Vid [190] lifts a satellite image to a street-view panoramic video, addressing the altitude gap between aerial mapping and ground-level perception. Wild-GS [171] solves the unconstrained internet-photo case by disentangling transient and persistent Gaussian appearance in 3DGS. InSpatio-WorldFM [203] is an open-source real-time generative frame model targeting spatial world understanding at scale.

8.3 Generative World Models and Foundation Models

World foundation models represent the apex of the NVS-to-world-modeling trajectory: systems that jointly handle viewpoint, time, and action, trained at internet scale. We identify three design axes that differentiate current approaches.

Sequence vs. spatial representation. GAIA-1 [117] treats driving scenes as sequences of video, text, and action tokens, training an autoregressive transformer to predict the next token. This approach achieves strong temporal coherence but lacks an explicit 3-D representation, limiting viewpoint flexibility. Genie [124] learns a latent action space from unsupervised internet videos, enabling interactive game-world generation without action labels—an important step toward self-supervised world model pretraining.

3-D-grounded generation. Genie 2 [125] introduces persistent 3-D structure into the Genie framework, generating scenes that maintain geometry across viewpoints and support first-person interactive navigation—the closest current system to an NVS-grounded world model for open-ended environments. WonderWorld [129] takes a complementary approach: instead of learning from large-scale video, it uses instant 3DGS initialisation from a single image combined with diffusion inpainting to generate interactive 3-D scenes on the fly. WonderJourney [128] chains text-guided scene generation steps to produce perpetual explorable 3-D environments.

Physical-world foundation models. Cosmos [126] is the most ambitious current effort: a family of world foundation models pretrained on internet-scale video with explicit physical world conditioning, designed for downstream fine-tuning to robotics and autonomous driving. Pandora [127] interleaves language and video tokens, enabling language-driven world simulation with both visual and linguistic output.

Table 7 compares world modeling methods.

8.4 Embodied AI and Robotics

NVS is more than a rendering tool in robotics: it provides the *spatial prior* that allows agents to reason about unseen parts of the environment before taking action. Three use patterns have emerged.

Pose estimation and object localisation. iNeRF [131] inverts a trained NeRF to estimate 6-DoF camera pose by gradient descent on the rendered photometric error—the first demonstration that NeRF's differentiable renderer can be repurposed for inference as well as reconstruction. Dex-NeRF [130] applies NeRF to transparent-object grasping, where conventional depth sensors fail, by

Table 8: NVS vs. World Modeling capability boundary.

Capability	NVS (static)	World Model
Viewpoint-consistent rendering (D1)	✓	✓
Sparse/single-image input (D2)	∼	✓
Temporal dynamics (D3)	–	✓
Action/agent conditioning (D4)	–	✓
Physical plausibility (D5)	–	✓
Cross-scene generalisation	∼	✓
Causal scene understanding	–	✓

leveraging NeRF's ability to model view-dependent appearance to infer geometry from specular reflections.

Visual planning and manipulation. NeRF-Supervisor [132] uses a pre-trained NeRF to synthesise robot-wrist camera views at hypothetical future poses, enabling visual planning without physical rollouts. GaussianGrasping [133] replaces NeRF with 3DGS for real-time manipulation planning, exploiting the explicit Gaussian representation for fast collision checking and grasp pose optimisation.

Simulation-to-real transfer. RoboGSim [134] reconstructs robot workspaces in 3DGS and renders synthetic training demonstrations from novel viewpoints, directly closing the sim-to-real gap for visuomotor policy learning.

8.5 Human Novel View Synthesis

Synthesising novel views of humans is a critical subproblem with applications in AR/VR, telepresence, and embodied AI.

Neural implicit human models. Neural Body [189] anchors structured latent codes on a parametric human body (SMPL), enabling novel view and pose synthesis. Rotationally-Temporally Consistent [188] enforces both rotation and temporal consistency for human performance video synthesis. NHR [150] adopts multi-view point-cloud-based rendering for dynamic human synthesis.

Generalizable and diffusion-based human NVS. CFSynthesis [178] enables controllable and free-view 3-D human video synthesis via a disentangled representation. iVS-Net [181] learns human view synthesis from internet videos without any 3-D supervision. Blur2Sharp [187] addresses human novel pose and view synthesis from blurry input with a generative prior. Diffuman4D [186] achieves 4-D consistent human view synthesis from sparse-view videos via video diffusion. EVA-Gaussian [185] achieves real-time 3DGS-based human NVS under diverse multi-view settings.

8.6 The Boundary: NVS vs. World Modeling

Table 8 delineates the capability boundary.

9 Datasets and Benchmarks

9.1 Object-Level Datasets

ShapeNet [145]: 51K 3-D CAD models; foundation for single-image reconstruction training. **Objaverse** [146]: 800K+ 3-D objects; enables LRM-scale training. **Objaverse-XL** [147]: 10M+ 3-D assets. **Co3D** [148]: 19K+ annotated per-category real video sequences for category-level NVS. **MVImgNet** [149]: 238K multi-view real images across 238 categories.

9.2 Static Scene Datasets

NeRF-Synthetic [17]: 8 synthetic bounded objects. **LLFF** [139]: 8 forward-facing real scenes. **DTU** [140]: 124 controlled MVS scenes. **Tanks and Temples** [141]: large-scale indoor/outdoor. **Mip-NeRF 360** [20]: 9 unbounded 360° scenes. **ScanNet** [142]: 1513 indoor RGB-D sequences.

Table 9: Key NVS datasets. RT=real, Syn=synthetic.

Dataset	Year	Type	Dyn.	GT Depth	Scale	Primary Use
ShapeNet [145]	2015	Syn	–	✓	51K obj.	Recon. training
DTU [140]	2014	RT	–	✓	124 scn.	MVS/NVS
LLFF [139]	2019	RT	–	–	8 scn.	Fwd-facing
NeRF-Synth. [17]	2020	Syn	–	✓	8 obj.	NeRF eval.
RealEstate10K [144]	2018	RT	–	–	10K vid.	Generalizable NVS
Tanks&Temples [141]	2017	RT	–	✓	Large	Reconstruction
D-NeRF [37]	2021	Syn	✓	–	6 obj.	Dynamic NeRF
HyperNeRF [40]	2021	RT	✓	–	Small	Topol. change
Mip-NeRF360 [20]	2022	RT	–	–	9 scn.	Unbounded
Neu3D [41]	2022	RT	✓	–	Human	Dyn. 3DGS
Co3D [148]	2021	RT	∼	–	19K seq.	Category NVS
ScanNet++ [143]	2023	RT	–	✓	Indoor	Generalise
Objaverse [146]	2023	Syn	∼	✓	800K obj.	LRM training
Waymo Open [153]	2020	RT	✓	✓	Urban	Driving sim.
nuScenes [154]	2020	RT	✓	∼	Urban	World model

Table 10: NVS quality on NeRF-Synthetic (Blender) and LLFF datasets. Higher PSNR/SSIM and lower LPIPS are better.

Method	NeRF-Synthetic			LLFF		
	PSNR↑	SSIM↑	LPIPS↓	PSNR↑	SSIM↑	LPIPS↓
NeRF [17]	31.01	0.947	0.081	26.50	0.811	0.250
Mip-NeRF [19]	33.09	0.961	0.043	26.70	0.829	0.177
Instant-NGP [23]	33.18	0.963	0.037	25.59	0.817	0.224
TensoRF [24]	33.14	0.963	0.047	26.73	0.839	0.204
Mip-NeRF360 [20]	–	–	–	29.40	0.901	0.127
3DGS [46]	33.32	0.963	0.044	27.21	0.815	0.183
Mip-Splatting [47]	34.66	0.973	0.019	27.79	0.827	0.163
Zip-NeRF [21]	–	–	–	30.38	0.921	0.100

ScanNet++ [143]: DSLR + iPhone high-resolution indoor. **RealEstate10K** [144]: 10K real estate videos for generalizable NVS training.

9.3 Dynamic Scene Datasets

D-NeRF [37]: 6 monocular synthetic dynamic objects. **HyperNeRF** [40]: monocular topological changes. **NHR** [150]: multi-view human performance. **CMU Panoptic** [151]: dense multi-view social activities. **Neu3D** [41]: multi-view dynamic human sequences. **iPhone Dataset** [152]: casually-captured phone videos.

9.4 Autonomous Driving and Large-Scale

Waymo Open [153]: multi-camera + LiDAR driving. **nuScenes** [154]: 360° driving + HD maps. **KITTI-360** [155]: urban driving with panoptic labels. **Argoverse 2** [156]: high-resolution 360° driving. **PandaSet** [157]: urban LiDAR+camera dataset.

Table 9 provides a structured comparison of key datasets.

10 Evaluation Metrics and Protocols

10.1 Standard Image Quality Metrics

PSNR (peak signal-to-noise ratio): pixel-level fidelity. **SSIM** [158]: structural similarity. **LPIPS** [159]: perceptual similarity via deep features; best correlated with human judgement.

Table 10 shows representative benchmark results on NeRF-Synthetic and LLFF.

10.2 Video and World Modelling Metrics

FVD [160] (Fréchet Video Distance): distributional quality of generated video sequences. **FID** [161]: Fréchet Inception Distance for image quality. **Temporal consistency**: average cosine similarity of

DINO features between consecutive frames. **Task-specific**: for driving, closed-loop collision rate, route completion, imitation gap; for robotics, task success rate.

10.3 Limitations of Current Evaluation

PSNR/SSIM penalise perceptually plausible but pixel-misaligned generations, especially under large viewpoint changes. FVD captures perceptual quality but not geometric accuracy. No existing benchmark jointly evaluates geometry, dynamics, action responsiveness, and physical plausibility. This evaluation gap is itself a major open problem.

Towards systematic evaluation. The work on Systematic Evaluation of NVS [201] provides a principled analysis of how existing benchmarks differ and proposes controlled evaluation protocols separating appearance and geometry quality. Non-Aligned Reference IQA [202] addresses the fundamental problem of quality assessment when reference and synthesised views are not pixel-aligned, proposing a learned non-aligned metric for NVS. VGGT-X [165] extends VGGT to dense NVS, providing strong baselines for the emerging class of feed-forward 3-D foundation models.

11 Open Problems and Research Roadmap

11.1 The Geometry–Generation Trade-Off

Reconstruction methods deliver geometric precision but lack generative priors; diffusion methods provide rich priors but can violate multi-view consistency. Methods like ReconFusion, Cat3D, and ViewCrafter begin to bridge this gap, but the fundamental tension remains. Principled solutions (*e.g.*, geometry-grounded diffusion, diffusion-seeded 3DGS) are an active frontier.

11.2 Long-Horizon Temporal Coherence

Existing 4-D methods maintain coherence over a few seconds. Two promising directions: (1) *Persistent scene memory*: Gaussian maps that accumulate observations across time and preserve global consistency; (2) *World model pre-training*: leveraging internet-scale video prediction as a temporal prior that can be fine-tuned to specific domains.

11.3 Physical Grounding

Current methods learn appearance correlations, not physical laws. Encouraging directions include: PAC-NeRF [138] for continuum mechanics integration; physics-informed deformable Gaussians for material simulation. The central challenge is integrating differentiable physics into the rendering pipeline in a way that scales to complex, open-world scenes.

11.4 Open-Vocabulary Semantic Integration

Methods like LERF [135] ground CLIP features in 3-D via NeRF, enabling open-vocabulary 3-D queries. LangSplat [136] brings language grounding to 3DGS. Feature3DGS [137] distils various 2-D foundation model features. Extending these to dynamic, actionable scenes is a key step toward language-grounded world models.

11.5 Evaluation and Benchmarks

We propose a five-axis world-modelling benchmark: (1) 3-D consistency under adversarial viewpoints; (2) temporal coherence over long horizons ($>$30 s); (3) action responsiveness in interactive settings; (4) physical plausibility via physics-engine verification; (5) open-world generalisation from a single image.

11.6 Research Roadmap

R1. Unified architectures combining explicit 3-D with diffusion priors.

R2. Internet-scale feed-forward generalisation for open-world NVS.

R3. Differentiable physics priors in 4-D world models.

R4. Language and semantic grounding for action-conditioned generation.

R5. Comprehensive world-modelling evaluation suite.

12 Conclusion

We have surveyed Novel View Synthesis through the lens of Generative World Modeling, tracing the evolution from classical image-based rendering through NeRF, 3D Gaussian Splatting, feed-forward reconstruction, and diffusion-based synthesis toward 4-D world generation. Our unified continuous-spectrum taxonomy organises more than 160 methods across three axes (input regime, temporal scope, controllability) and makes explicit their world-modelling implications.

The field is at an inflection point. Spatial grounding from NVS and temporal/action capabilities from world modeling are actively converging, as evidenced by systems such as Genie 2, Cosmos, ViewCrafter, and the neural driving simulators. Closing the geometry–generation trade-off, achieving long-horizon coherence, and grounding predictions in physical laws are the primary challenges ahead.

Draft v0.3—ongoing work. Content will be further expanded in subsequent versions.

References

[1] M. Levoy and P. Hanrahan. Light field rendering. *SIGGRAPH*, 1996.

[2] L. McMillan and G. Bishop. Plenoptic modeling. *SIGGRAPH*, 1995.

[3] S. E. Chen and L. Williams. View interpolation for image synthesis. *SIGGRAPH*, 1993.

[4] P. E. Debevec, C. J. Taylor, and J. Malik. Modeling and rendering architecture from photographs. *SIGGRAPH*, 1996.

[5] S. J. Gortler, R. Grzeszczuk, R. Szeliski, and M. F. Cohen. The lumigraph. *SIGGRAPH*, 1996.

[6] E. H. Adelson and J. R. Bergen. The plenoptic function and the elements of early vision. In *Computational Models of Visual Processing*, 1991.

[7] J. L. Schönberger and J.-M. Frahm. Structure-from-motion revisited. *CVPR*, 2016.

[8] J. L. Schönberger, E. Zheng, J.-M. Frahm, and M. Pollefeys. Pixelwise view selection for unstructured multi-view stereo. *ECCV*, 2016.

[9] A. Tewari *et al.*Advances in neural rendering. *Computer Graphics Forum (CGF)*, 41(2):703–735, 2022.

[10] K. Gao *et al.*NeRF: Neural radiance field in 3D vision—a comprehensive review. *arXiv:2210.00379*, 2022.

[11] T. Zhu *et al.*A survey on image-based rendering. *arXiv*, 2023.

[12] T. Wu *et al.*Recent advances in 3D Gaussian Splatting. *arXiv:2403.11134*, 2024.

[13] B. Fei *et al.*3D Gaussian Splatting as a new era of novel view synthesis. *IEEE TVCG*, 2024.

[14] K. Chen *et al.*A survey of novel view synthesis from autonomous driving perspective. *arXiv*, 2024.

[15] J. Zhang *et al.*Advances in feed-forward 3D reconstruction and view synthesis: A survey. *arXiv*, 2025.

[16] S. He *et al.*A survey on 3D Gaussian Splatting applications. *arXiv:2308.04079*, 2025.

[17] B. Mildenhall, P. P. Srinivasan, M. Tancik, J. T. Barron, R. Ramamoorthi, and R. Ng. NeRF: Representing scenes as neural radiance fields for view synthesis. *ECCV*, 2020.

[18] R. Martin-Brualla, N. Radwan, M. S. M. Sajjadi, J. T. Barron, A. Dosovitskiy, and D. Duckworth. NeRF in the Wild: Neural radiance fields for unconstrained photo collections. *CVPR*, 2021.

[19] J. T. Barron, B. Mildenhall, M. Tancik, P. Hedman, R. Martin-Brualla, and P. P. Srinivasan. Mip-NeRF: A multiscale representation for anti-aliasing neural radiance fields. *ICCV*, 2021.

[20] J. T. Barron, B. Mildenhall, D. Verbin, P. P. Srinivasan, and P. Hedman. Mip-NeRF 360: Unbounded anti-aliased neural radiance fields. *CVPR*, 2022.

[21] J. T. Barron, B. Mildenhall, D. Verbin, P. P. Srinivasan, and P. Hedman. Zip-NeRF: Anti-aliased grid-based neural radiance fields. *ICCV*, 2023.

[22] D. Verbin, P. Hedman, B. Mildenhall, T. Zickler, J. T. Barron, and P. P. Srinivasan. Ref-NeRF: Structured view-dependent appearance for neural radiance fields. *CVPR*, 2022.

[23] T. Müller, A. Evans, C. Schied, and A. Keller. Instant neural graphics primitives with a multiresolution hash encoding. *ACM TOG*, 2022.

[24] A. Chen, Z. Xu, A. Geiger, J. Yu, and H. Su. TensoRF: Tensorial radiance fields. *ECCV*, 2022.

[25] S. Fridovich-Keil, A. Yu, M. Tancik, Q. Chen, B. Recht, and A. Kanazawa. Plenoxels: Radiance fields without neural networks. *CVPR*, 2022.

[26] C. Sun, M. Sun, and H.-T. Chen. Direct voxel grid optimization: Super-fast convergence for radiance fields reconstruction. *CVPR*, 2022.

[27] P. Wang, L. Liu, Y. Liu, C. Theobalt, T. Komura, and W. Wang. NeuS: Learning neural implicit surfaces by volume rendering for multi-view reconstruction. *NeurIPS*, 2021.

[28] L. Yariv, J. Gu, Y. Kasten, and Y. Lipman. Volume rendering of neural implicit surfaces. *NeurIPS*, 2021.

[29] L. Yariv, P. Hedman, C. Reiser, D. Verbin, P. P. Srinivasan, R. Bastian-Smith, J. T. Barron, and R. Martin-Brualla. BakedSDF: Meshing neural SDFs for real-time view synthesis. *SIGGRAPH*, 2023.

[30] A. Yu, V. Ye, M. Tancik, and A. Kanazawa. pixelNeRF: Neural radiance fields from one or few images. *CVPR*, 2021.

[31] Q. Wang, Z. Wang, K. Genova, P. P. Srinivasan, H. Zhou, J. T. Barron, R. Martin-Brualla, N. Snavely, and T. Funkhouser. IBRNet: Learning multi-view image-based rendering. *CVPR*, 2021.

[32] A. Chen, Z. Xu, F. Zhao, X. Zhang, F. Xiang, J. Yu, and H. Su. MVSNeRF: Fast generalizable radiance field reconstruction from multi-view stereo. *ICCV*, 2021.

[33] S. Wang, T. Leimkühler, and C. Theobalt. Generalizable novel view synthesis via neural radiance field transformers. *arXiv*, 2022.

[34] M. S. M. Sajjadi, H. Meyer, E. Pot, U. Bergmann, K. Greff, N. Radwan, S. Vora, and M. Hessel. Scene representation transformer: Geometry-free novel view synthesis through set-latent scene representations. *CVPR*, 2022.

[35] J. Gu *et al.*NerfDiff: Single-image view synthesis with nerf-guided distillation. *ICML*, 2023.

[36] R. Shi, X. Wei, C. Wang, and H. Su. ZeroRF: Fast sparse view 360° reconstruction with zero pretraining. *arXiv:2312.09967*, 2023.

[37] A. Pumarola, E. Corona, G. Pons-Moll, and F. Moreno-Noguer. D-NeRF: Neural radiance fields for dynamic scenes. *CVPR*, 2021.

[38] Z. Li, S. Niklaus, N. Snavely, and O. Wang. Neural scene flow fields for space-time view synthesis of dynamic scenes. *CVPR*, 2021.

[39] K. Park, U. Sinha, J. T. Barron, S. Bouaziz, D. B. Goldman, S. M. Seitz, and R. Martin-Brualla. Nerfies: Deformable neural radiance fields. *ICCV*, 2021.

[40] K. Park, U. Sinha, P. Hedman, J. T. Barron, S. Bouaziz, D. B. Goldman, R. Martin-Brualla, and S. M. Seitz. HyperNeRF: A higher-dimensional representation for topologically varying neural radiance fields. *ACM TOG*, 2021.

[41] Z. Li, S. Niklaus, N. Snavely, and O. Wang. Neural 3D video synthesis from multi-view video. *CVPR*, 2022.

[42] J. Fang, T. Yi, X. Wang, L. Xie, X. Zhang, W. Liu, M. Niëßner, and L. Tian. Fast dynamic radiance fields with time-aware neural voxels. *SIGGRAPH Asia*, 2022.

[43] S. Fridovich-Keil, G. Meanti, F. R. Warburg, B. Recht, and A. Kanazawa. K-Planes: Explicit radiance fields in space, time, and appearance. *CVPR*, 2023.

[44] A. Cao and J. Johnson. HexPlane: A fast representation for dynamic scenes. *CVPR*, 2023.

[45] L. Song *et al.*NeRFPlayer: A streamable dynamic scene representation with decomposed neural radiance fields. *IEEE TVCG*, 2023.

[46] B. Kerbl, G. Kopanas, T. Leimkühler, and G. Drettakis. 3D Gaussian Splatting for real-time radiance field rendering. *ACM TOG*, 2023.

[47] Z. Yu, A. Chen, B. Huang, T. Sattler, and A. Geiger. Mip-Splatting: Alias-free 3D Gaussian Splatting. *CVPR*, 2024.

[48] Z. Ye, Z. Liu, W. Song, L. Liu, X. Zhao, T. Zhu, and B. Zhang. AbsGS: Recovering fine details for 3D Gaussian Splatting. *NeurIPS*, 2024.

[49] T. Lu *et al.*Scaffold-GS: Structured 3D Gaussians for view-adaptive rendering. *CVPR*, 2024.

[50] Z. Fan, K. Wang, K. Wen, Z. Zhu, D. Xu, and Z. Wang. LightGaussian: Unbounded 3D Gaussian compression with 15× reduction. *NeurIPS*, 2023.

[51] J. Lee *et al.*Compact 3D Gaussian representation for radiance field. *arXiv:2311.13681*, 2024.

[52] S. Girish, K. Gupta, and A. Shrivastava. EAGLES: Efficient accelerated 3D Gaussians with lightweight encodings. *ECCV*, 2024.

[53] Y. Chen, L. Wu, C. Gao, J. Li, and C. Liu. HAC: Hash-grid assisted context for 3D Gaussian Splatting compression. *ECCV*, 2024.

[54] A. Guédon and V. Lepetit. SuGaR: Surface-aligned Gaussian Splatting for efficient 3D mesh reconstruction and high-quality mesh rendering. *CVPR*, 2024.

[55] B.-S. Huang, Z. Yu, A. Chen, A. Geiger, and S. Gao. 2D Gaussian Splatting for geometrically accurate radiance fields. *SIGGRAPH*, 2024.

[56] Z. Yu, T. Sattler, and A. Geiger. Gaussian Opacity Fields: Efficient adaptive surface reconstruction in unbounded scenes. *arXiv:2404.10772*, 2024.

[57] Y. Chen, Z. Chen, C. Zhang, F. Wang, X. Yang, Y. Wang, Z. Cai, L. Yang, H. Liu, and G. Lin. GaussianEditor: Swift and controllable 3D editing with Gaussian Splatting. *CVPR*, 2024.

[58] M. Ye, M. Danelljan, F. Yu, and L. Ke. Gaussian Grouping: Segment and edit anything in 3D scenes. *ECCV*, 2024.

[59] X. Zhou *et al.*GALA3D: Towards text-to-3D complex scene generation via layout-guided generative Gaussian Splatting. *ICML*, 2024.

[60] G. Wu, T. Yi, J. Fang, L. Xie, X. Zhang, W. Wei, W. Liu, Q. Tian, and M. Wang. 4D Gaussian Splatting for real-time dynamic scene rendering. *CVPR*, 2024.

[61] Z. Yang, X. Gao, Y. Sun, Y. Huang, X. Lyu, W. Zhou, X. Jiao, X. Qi, Y. Jin, D. Yang, X. Duan, F. Yu, L. Liu, P. Zhao, and A. Qu. Deformable 3D Gaussians for high-fidelity monocular dynamic reconstruction. *CVPR*, 2024.

[62] Z. Yang, X. Gao, W. Zhou, S. Jiao, Y. Zhang, and X. Jin. Real-time photorealistic dynamic scene representation and rendering with 4D Gaussian Splatting. *ICLR*, 2024.

[63] Y.-H. Huang, Y.-T. Sun, Z. Yang, X. Lyu, Y.-P. Cao, and X. Qi. SC-GS: Sparse-controlled Gaussian Splatting for editable dynamic scenes. *CVPR*, 2024.

[64] Q. Lin *et al.*Gaussian-Flow: 4D reconstruction with dynamic 3D Gaussian particle. *CVPR*, 2024.

[65] Q. Wang *et al.*Shape of Motion: 4D reconstruction from a single video. *arXiv:2407.13764*, 2024.

[66] D. Charatan, S. Li, A. Tagliasacchi, and V. Sitzmann. pixelSplat: 3D Gaussian splats from image pairs for scalable generalizable 3D reconstruction. *CVPR*, 2024.

[67] Y. Chen *et al.*MVSplat: Efficient 3D Gaussian Splatting from sparse multi-view images. *ECCV*, 2024.

[68] C. Zhang *et al.*TranSplat: Generalizable 3D Gaussian Splatting from sparse multi-view images with transformers. *arXiv:2408.13912*, 2024.

[69] Z. Zhu, Z. Fan, Y. Jiang, and Z. Wang. FSGS: Real-time few-shot view synthesis using Gaussian Splatting. *CVPR*, 2024.

[70] H. Xu *et al.*DepthSplat: Connecting Gaussian Splatting and depth. *arXiv*, 2024.

[71] B. Smart *et al.*Splatt3R: Zero-shot Gaussian splatting from uncalibrated image pairs. *arXiv:2408.13912*, 2024.

[72] Z.-X. Zou, Z. Yu, Y.-B. Guo, Y. Li, L. Liang, C.-J. Cao, and Y. Zhang. Triplane meets Gaussian Splatting: Fast and generalizable single-view 3D reconstruction with transformers. *CVPR*, 2024.

[73] S. Wang *et al.*DUSt3R: Geometric 3D vision made easy. *CVPR*, 2024.

[74] V. Leroy *et al.*MASt3R: Grounding Image Matching in 3D with MASt3R. *ECCV*, 2024.

[75] J. Zhang *et al.*MonST3R: A simple approach for estimating geometry in the presence of motion. *arXiv:2410.03825*, 2024.

[76] H. Wang *et al.*VGGT: Visual geometry grounded deep structured tokens. *arXiv:2503.11651*, 2025.

[77] Y. Hong, K. Zhang, J. Gu, S. Bi, Y. Zhou, D. Liu, F. Liu, K. Sunkavalli, T. Bui, and H. Tan. LRM: Large reconstruction model for single image to 3D. *ICLR*, 2024.

[78] J. Li *et al.*Instant3D: Fast text-to-3D with sparse-view generation and large reconstruction model. *ICLR*, 2024.

[79] Z. He and H. Wang. OpenLRM: Open-source large reconstruction models. *arXiv:2311.04400*, 2023.

[80] J. Tang *et al.*LGM: Large multi-view Gaussian model for high-resolution 3D content creation. *ECCV*, 2024.

[81] Z. Wang *et al.*CRM: Single image to 3D textured mesh with convolutional reconstruction model. *NeurIPS*, 2024.

[82] Y. Yao, Z. Luo, S. Li, T. Fang, and L. Quan. MVSNet: Depth inference for unstructured multi-view stereo. *ECCV*, 2018.

[83] F. Wang *et al.*IterMVS: Iterative probability estimation for efficient multi-view stereo. *CVPR*, 2022.

[84] R. Peng *et al.*Rethinking depth estimation for multi-view stereo: A unified representation. *CVPR*, 2022.

[85] K. Schwarz, Y. Liao, M. Niemeyer, and A. Geiger. GRAF: Generative radiance fields for 3D-aware image synthesis. *NeurIPS*, 2020.

[86] E. R. Chan, M. Monteiro, P. Kellnhofer, J. Wu, and G. Wetzstein. pi-GAN: Periodic implicit generative adversarial networks for 3D-aware image synthesis. *CVPR*, 2021.

[87] M. Niemeyer and A. Geiger. GIRAFFE: Representing scenes as compositional generative neural feature fields. *CVPR*, 2021.

[88] E. R. Chan *et al.*Efficient geometry-aware 3D generative adversarial networks. *CVPR*, 2022.

[89] J. Gu, L. Liu, P. Wang, and C. Theobalt. StyleNeRF: A style-based 3D-aware generator for high-resolution image synthesis. *ICLR*, 2022.

[90] R. Rombach, A. Blattmann, D. Lorenz, P. Esser, and B. Ommer. High-resolution image synthesis with latent diffusion models. *CVPR*, 2022.

[91] R. Liu, R. Wu, B. V. Hoorick, P. Tokmakov, S. Zakharov, and C. Vondrick. Zero-1-to-3: Zero-shot one image to 3D object. *ICCV*, 2023.

[92] Y. Liu, C. Lin, Z. Zeng, X. Long, L. Liu, T. Komura, and W. Wang. SyncDreamer: Generating multiview-consistent images from a single-view image. *arXiv:2309.03453*, 2023.

[93] M. Liu *et al.*One-2-3-45: Any single image to 3D mesh in 45 seconds without per-shape optimization. *NeurIPS*, 2023.

[94] R. Shi *et al.*Zero123++: A single image to consistent multi-view diffusion base model. *arXiv:2310.15110*, 2023.

[95] S. Tang, F. Zhang, J. Chen, P. Wang, and Y. Furukawa. MVDiffusion: Enabling holistic multi-view image generation with correspondence-aware diffusion. *NeurIPS*, 2023.

[96] H. Long *et al.*Wonder3D: Single image to 3D using cross-domain diffusion. *CVPR*, 2024.

[97] X. Kong *et al.*EscherNet: A generative model for scalable view synthesis. *CVPR*, 2024.

[98] K. Sargent *et al.*ZeroNVS: Zero-shot 360-degree view synthesis from a single real image. *CVPR*, 2024.

[99] C. Zheng *et al.*Free3D: Consistent novel view synthesis without 3D representation. *ICLR*, 2024.

[100] E. R. Chan *et al.*Generative novel view synthesis with 3D-aware diffusion models. *ICCV*, 2023.

[101] Y. Shi *et al.*MVDream: Multi-view diffusion for 3D generation. *ICLR*, 2024.

[102] B. Poole, A. Jain, J. T. Barron, and B. Mildenhall. DreamFusion: Text-to-3D using 2D diffusion. *ICLR*, 2023.

[103] C.-H. Lin, J. Gao, L. Tang, T. Takikawa, X. Zeng, X. Huang, K. Kreis, S. Fidler, M.-Y. Liu, and T.-Y. Lin. Magic3D: High-resolution text-to-3D content creation. *CVPR*, 2023.

[104] Z. Wang *et al.*ProlificDreamer: High-fidelity and diverse text-to-3D generation with variational score distillation. *NeurIPS*, 2023.

[105] R. Chen *et al.*Fantasia3D: Disentangling geometry and appearance for high-quality text-to-3D content creation. *ICCV*, 2023.

[106] T. Yi *et al.*GaussianDreamer: Fast generation from text to 3D Gaussians by bridging 2D and 3D diffusion models. *CVPR*, 2024.

[107] J. Chung *et al.*LucidDreamer: Domain-free generation of 3D Gaussian Splatting scenes. *arXiv:2311.13384*, 2023.

[108] R. Gao *et al.*CAT3D: Create anything in 3D with multi-view diffusion models. *arXiv:2405.10314*, 2024.

[109] R. Wu *et al.*ReconFusion: 3D reconstruction with diffusion priors. *CVPR*, 2024.

[110] W. Yu *et al.*ViewCrafter: Taming video diffusion models for high-fidelity novel view synthesis. *arXiv:2409.02048*, 2024.

[111] D. Watson *et al.*4DiM: 4D diffusion for dynamic novel view synthesis with global-local attention. *arXiv:2311.17905*, 2024.

[112] Y. Xie *et al.*SV4D: Dynamic 3D content generation with multi-frame and multi-view consistency. *arXiv:2407.09787*, 2024.

[113] Y. Jiang *et al.*Consistent4D: Consistent 360° dynamic object generation from monocular video. *ICLR*, 2024.

[114] Y. Ren *et al.*DreamGaussian4D: Generative 4D Gaussian Splatting. *arXiv:2312.17142*, 2023.

[115] D. Ha and J. Schmidhuber. World models. *arXiv:1803.10122*, 2018.

[116] Y. LeCun. A path towards autonomous machine intelligence. *OpenReview*, 2022.

[117] A. Hu *et al.*GAIA-1: A generative world model for autonomous driving. *arXiv:2309.17080*, 2023.

[118] Z. Yang, Y. Chen, J. Wang, S. Manivasagam, W.-C. Ma, A. K. Yang, and R. Urtasun. UniSim: A neural closed-loop sensor simulator. *CVPR*, 2023.

[119] X. Wu *et al.*MARS: An instance-aware, modular and realistic simulator for autonomous driving. *CICAI*, 2023.

[120] Z. Chen *et al.*OmniRe: Omni urban scene reconstruction. *arXiv:2408.16760*, 2024.

[121] J. Yang *et al.*EmerNeRF: Emergent spatial-temporal scene decomposition via self-supervision. *ICLR*, 2024.

[122] Y. Yan *et al.*Street Gaussians: Modelling dynamic urban scenes with Gaussian Splatting. *ECCV*, 2024.

[123] Y. Chen *et al.*Periodic vibration Gaussian: Dynamic urban scene reconstruction and real-time rendering. *arXiv:2311.18561*, 2023.

[124] J. Bruce *et al.*Genie: Generative interactive environments. *ICML*, 2024.

[125] J. Parker-Holder *et al.*Genie 2: A large-scale foundation world model. *Google DeepMind Technical Report*, 2024.

[126] N. Agarwal *et al.*Cosmos world foundation model technical report. *arXiv:2501.03575*, 2025.

[127] J. Xiang *et al.*Pandora: Towards general world model with natural language actions and video states. *arXiv:2406.09455*, 2024.

[128] H. Yu *et al.*WonderJourney: Going from anywhere to everywhere. *CVPR*, 2024.

[129] H. Yu *et al.*WonderWorld: Interactive 3D scene generation using a single image. *arXiv:2406.09379*, 2024.

[130] J. Ichnowski, Y. Avigal, J. Kerr, and K. Goldberg. Dex-NeRF: Using a neural radiance field to grasp transparent objects. *CoRL*, 2021.

[131] L.-Y. Yen-Chen, P. Florence, J. T. Barron, A. Rodriguez, P. Isola, and T.-Y. Lin. iNeRF: Inverting neural radiance fields for pose estimation. *IROS*, 2021.

[132] Y. Ze *et al.*NeRF-Supervisor: Learning dense robot manipulation policies with neural radiance fields. *ICRA*, 2023.

[133] O. Shorinwa *et al.*Splat-MOVER: Multi-stage, open-vocabulary robotic manipulation via editable Gaussian splatting. *arXiv*, 2024.

[134] K. Li *et al.*RoboGSim: A real2sim2real robotic Gaussian splatting simulator. *arXiv:2411.11839*, 2024.

[135] J. Kerr, C. M. Kim, K. Goldberg, A. Kanazawa, and M. Tancik. LERF: Language embedded radiance fields. *ICCV*, 2023.

[136] M. Qin *et al.*LangSplat: 3D language Gaussian Splatting. *CVPR*, 2024.

[137] S. Zhou *et al.*Feature 3DGS: Supercharging 3D Gaussian Splatting to enable distilled feature fields. *CVPR*, 2024.

[138] Y. Li *et al.*PAC-NeRF: Physics augmented continuum neural radiance fields for geometry-agnostic system identification. *ICLR*, 2023.

[139] B. Mildenhall, P. P. Srinivasan, R. Ortiz-Cayon, N. K. Kalantari, R. Ramamoorthi, R. Ng, and A. Kar. Local light field fusion: Practical view synthesis with prescriptive sampling guidelines. *ACM TOG*, 2019.

[140] R. Jensen, A. Dahl, G. Vogiatzis, E. Tola, and H. Aanæs. Large scale multi-view stereopsis evaluation. *CVPR*, 2014.

[141] A. Knapitsch, J. Park, Q.-Y. Zhou, and V. Koltun. Tanks and temples: Benchmarking large-scale scene reconstruction. *ACM TOG*, 2017.

[142] A. Dai, A. X. Chang, M. Savva, M. Halber, T. Funkhouser, and M. Nießner. ScanNet: Richly-annotated 3D reconstructions of indoor scenes. *CVPR*, 2017.

[143] C. Yeshwanth, Y.-C. Liu, M. Nießner, and A. Dai. ScanNet++: A high-fidelity dataset of 3D indoor scenes. *ICCV*, 2023.

[144] T. Zhou, R. Tucker, J. Flynn, G. Fyffe, and N. Snavely. Stereo magnification: Learning view synthesis using multiplane images. *SIGGRAPH*, 2018.

[145] A. X. Chang *et al.*ShapeNet: An information-rich 3D model repository. *arXiv:1512.03012*, 2015.

[146] M. Deitke *et al.*Objaverse: A universe of annotated 3D objects. *CVPR*, 2023.

[147] M. Deitke *et al.*Objaverse-XL: A universe of 10M+ 3D objects. *NeurIPS*, 2024.

[148] J. Reizenstein, R. Shapovalov, P. Henzler, L. Sbordone, P. Labatut, and D. Novotny. Common objects in 3D: Large-scale learning and evaluation of real-life 3D category reconstruction. *ICCV*, 2021.

[149] X. Yu *et al.*MVImgNet: A large-scale dataset of multi-view images. *CVPR*, 2023.

[150] M. Wu *et al.*Multi-view neural human rendering. *CVPR*, 2020.

[151] H. Joo *et al.*Panoptic studio: A massively multiview system for social motion capture. *ICCV*, 2015.

[152] X. Guo *et al.*Neural 3D scene reconstruction with the Manhattan-world assumption. *CVPR*, 2022.

[153] P. Sun *et al.*Scalability in perception for autonomous driving: Waymo open dataset. *CVPR*, 2020.

[154] H. Caesar *et al.*nuScenes: A multimodal dataset for autonomous driving. *CVPR*, 2020.

[155] Y. Liao, J. Xie, and A. Geiger. KITTI-360: A novel dataset and benchmarks for urban scene understanding. *IEEE TPAMI*, 2022.

[156] B. Wilson *et al.*Argoverse 2: Next generation datasets for self-driving perception and forecasting. *NeurIPS Datasets & Benchmarks*, 2023.

[157] C. Xiao *et al.*PandaSet: Advanced sensor suite dataset for autonomous driving. *ITSC*, 2021.

[158] Z. Wang, A. C. Bovik, H. R. Sheikh, and E. P. Simoncelli. Image quality assessment: From error visibility to structural similarity. *IEEE TIP*, 13(4):600–612, 2004.

[159] R. Zhang, P. Isola, A. A. Efros, E. Shechtman, and O. Wang. The unreasonable effectiveness of deep features as a perceptual metric. *CVPR*, 2018.

[160] T. Unterthiner, S. van Steenkiste, K. Kurach, R. Marinier, M. Michalski, and S. Gelly. Towards accurate generative models of video: A new metric & challenges. *arXiv:1812.01717*, 2019.

[161] M. Heusel, H. Ramsauer, T. Unterthiner, B. Nessler, and S. Hochreiter. GANs trained by a two time-scale update rule converge to a local Nash equilibrium. *NeurIPS*, 2017.

[162] L. Yan, R. He, Y. Sun, and J. Yu. FrameNeRF: A simple and efficient framework for few-shot novel view synthesis. *arXiv:2401.11734*, 2024.

[163] Z. Xu *et al.*Novel view synthesis using DDIM inversion. *arXiv:2401.08224*, 2024.

[164] W. Chen *et al.*Enhancing novel view synthesis from extremely sparse views with SfM-free 3D Gaussian Splatting framework. *arXiv:2408.00908*, 2024.

[165] C. Liu *et al.*VGGT-X: When VGGT meets dense novel view synthesis. *arXiv:2503.xxxxx*, 2025.

[166] X. Han *et al.*Projected representation conditioning for high-fidelity novel view synthesis. *arXiv*, 2024.

[167] Y. Tang *et al.*A single image and multimodality is all you need for novel view synthesis. *arXiv*, 2024.

[168] S. Chen *et al.*GeoNVS: Geometry grounded video diffusion for novel view synthesis. *arXiv*, 2024.

[169] M. Kwak *et al.*ViVid-1-to-3: Novel view synthesis with video diffusion models. *arXiv:2312.01305*, 2024.

[170] Y. Li *et al.*NVComposer: Boosting generative novel view synthesis with multiple sparse and unposed images. *arXiv:2409.11105*, 2024.

[171] Z. Xu *et al.*Wild-GS: Real-time novel view synthesis from unconstrained photo collections. *NeurIPS*, 2024.

[172] J. Shi *et al.*Novel view extrapolation with video diffusion priors. *arXiv*, 2024.

[173] X. Zhou *et al.*Recollection from Pensieve: Novel view synthesis via learning from uncalibrated videos. *arXiv*, 2024.

[174] Q. Wang *et al.*Mixed neural voxels for fast multi-view video synthesis. *ICCV*, 2023.

[175] J. Luiten *et al.*Dynamic 3D Gaussians: Tracking by persistent dynamic view synthesis. *3DV*, 2024.

[176] H. Kim *et al.*Self-calibrating 4D novel view synthesis from monocular videos using Gaussian Splatting. *arXiv*, 2024.

[177] J. Han *et al.*Reconstruct, inpaint, test-time finetune: Dynamic novel-view synthesis from monocular videos. *arXiv*, 2024.

[178] Z. Li *et al.*CFSynthesis: Controllable and free-view 3D human video synthesis. *arXiv*, 2023.

[179] Z. Wang *et al.*Decoupling dynamic monocular videos for dynamic view synthesis. *arXiv*, 2024.

[180] H. Li *et al.*Diffusion priors for dynamic view synthesis from monocular videos. *arXiv*, 2024.

[181] Y. Chen *et al.*iVS-Net: Learning human view synthesis from internet videos. *arXiv*, 2024.

[182] X. Li *et al.*Dynamic view synthesis from small camera motion videos. *arXiv*, 2024.

[183] Y. Du *et al.*Fast view synthesis of casual videos with Soup-of-Planes. *ICCV*, 2023.

[184] X. Zhao *et al.*Pseudo-generalized dynamic view synthesis from a video. *arXiv*, 2024.

[185] L. Hu *et al.*EVA-Gaussian: 3D Gaussian-based real-time human novel view synthesis under diverse multi-view settings. *arXiv*, 2024.

[186] Z. Zheng *et al.*Diffuman4D: 4D consistent human view synthesis from sparse-view videos with video diffusion models. *arXiv*, 2024.

[187] X. Dong *et al.*Blur2Sharp: Human novel pose and view synthesis with generative prior. *arXiv*, 2024.

[188] Y. Liu *et al.*Rotationally-temporally consistent novel view synthesis of human performance video. *CVPR*, 2021.

[189] S. Peng *et al.*Neural Body: Implicit neural representations with structured latent codes for novel view synthesis of dynamic humans. *CVPR*, 2021.

[190] S. Wu *et al.*Sat2Vid: Street-view panoramic video synthesis from a single satellite image. *ICCV*, 2021.

[191] Z. Liang *et al.*GauFRe: Gaussian deformation fields for real-time dynamic novel view synthesis. *arXiv:2312.11458*, 2024.

[192] Y. Sun *et al.*MoVieS: Motion-aware 4D dynamic view synthesis in one second. *arXiv*, 2024.

[193] Z. You *et al.*NVS-Solver: Video diffusion model as zero-shot novel view synthesizer. *arXiv:2405.15364*, 2024.

[194] Z. Li *et al.*Template-free articulated Gaussian Splatting for real-time reposable dynamic view synthesis. *NeurIPS*, 2024.

[195] J. Bao *et al.*Novel view synthesis from a few glimpses via test-time natural video completion. *arXiv*, 2024.

[196] E. Tretschk, A. Tewari, V. Golyanik, M. Zollhöfer, C. Lassner, and C. Theobalt. Non-rigid neural radiance fields: Reconstruction and novel view synthesis of a dynamic scene from monocular video. *ICCV*, 2021.

[197] Y.-C. Tsai *et al.*Consistent view synthesis with pose-guided diffusion models. *CVPR*, 2023.

[198] V. Voleti *et al.*SV3D: Novel multi-view synthesis and 3D generation from a single image using latent video diffusion. *ECCV*, 2024.

[199] R. Wu *et al.*CAT4D: Create anything in 4D with multi-view video diffusion models. *NeurIPS*, 2024.

[200] Y. Yan *et al.*StreetCrafter: Street view synthesis with controllable video diffusion models. *arXiv:2407.13758*, 2024.

[201] A. Chen *et al.*Systematic evaluation of novel view synthesis. *arXiv*, 2024.

[202] W. Zhang *et al.*Non-aligned reference image quality assessment for novel view synthesis. *arXiv*, 2024.

[203] K. Chen *et al.*InSpatio-WorldFM: An open-source real-time generative frame model. *arXiv*, 2025.

[204] J. T. Kajiya and B. P. Von Herzen. Ray tracing volume densities. *SIGGRAPH*, 1984.

[205] M. Tancik *et al.*Fourier features let networks learn high frequency functions in low dimensional domains. *NeurIPS*, 2020.

[206] J. Ho, A. Jain, and P. Abbeel. Denoising diffusion probabilistic models. *NeurIPS*, 2020.

[207] J. Song, C. Meng, and S. Ermon. Denoising diffusion implicit models. *ICLR*, 2021.

[208] T. Karras, S. Laine, M. Aittala, J. Hellsten, J. Lehtinen, and T. Aila. Analyzing and improving the image quality of StyleGAN. *CVPR*, 2020.

[209] V. Sitzmann, J. Martel, A. Bergman, D. Lindell, and G. Wetzstein. Implicit neural representations with periodic activation functions. *NeurIPS*, 2020.

[210] A. Kirillov *et al.*Segment Anything. *ICCV*, 2023.

[211] R. Liu *et al.*Zero-1-to-3-XL: Zero-shot one image to 3D object (extended scale). *arXiv:2311.09093*, 2023.

[212] P. Wang and Y. Shi. ImageDream: Image-prompt multi-view diffusion for 3D generation. *arXiv:2312.02201*, 2024.

[213] C. Wu *et al.*Direct3D: Scalable image-to-3D generation via 3D latent diffusion transformer. *arXiv:2405.14832*, 2024.